Paper Review: Fully Convolutional Networks for Semantic Segmentation (CVPR 2015)

Nithish karanam

1. Paper Summary

Full Convolutional Networks for Semantic Segmentation is a foundational paper that adapts deep classification networks into dense, pixel-wise prediction networks. Traditional CNNs have largely been built to aim towards image-level classification where the prediction is one label per image. The authors introduce a variation, though, where they replace the fully connected layers with convolutional layers so that the network is predicting a spatial map of predictions. This enables the network to be able to process an image of any dimension and generate a segmentation map that classifies every one of the pixels.
The novelty in the paper is the application of the skip connections and upsampling (or deconvolution) layers. Since convolutional layers will reduce the spatial size of the feature maps, the dimensions lost are restored with the help of the upsampling layers, and a clear segmentation map is formed. Skip connections from previous layers are also introduced, combining high-level context and fine details. This enables the output to have high-level context and accurate object boundary localization.
The other key aspect of the paper is the capability to train end-to-end. With the conversion of standard classification networks to full convolutional networks, it is possible to train the entire network within a single framework. This eliminates the necessity to have distinct, hand-tuned feature extraction stages or post-processing methods, and the model is more efficient and effective. Overall, the paper not only shows that deep learning is applicable to semantic segmentation, but establishes a new state of the art on segmentation on the PASCAL VOC benchmark. The paper spawned subsequent work in the field, with more advanced architectures built from these early ideas. This abstract summarizes how the paper shifts the paradigm from patch-based methods to an end-to-end trainable, global system for dense prediction tasks.

2. Experimental Results

The FCN paper's experimental part is strong, demonstrating the performance benefit of full convolutional networks on hard tasks including PASCAL VOC 2011/2012. The authors test several variants of their model—i.e., FCN-32s, FCN-16s, and FCN-8s—that differ in how they upsample the feature maps and introduce skip connections. FCN-32s uses a straightforwardly simple upsampling from the most bottom layer, resulting in less accurate segmentation maps. FCN-16s and FCN-8s, on the other hand, gradually combine features from preceding layers (via skip connections) to improve segmentation, with FCN-8s being the highest in resolution and performance.
The key measure in these experiments is mean Intersection over Union (mIoU), a measure of how much overlap there is between ground truth and segmentation prediction. From the results, we can see that FCN-32s performs decently, but by incorporating skip connections into FCN-16s and FCN-8s, there is vast mIoU improvement. There are qualitative results—visual segmentation maps illustrating how more refined object boundaries and details are recovered when these techniques are used.

Here is a sample table comparing the performance disparity between FCN versions (the below values are typical):

Model Variant	mIoU (%)	Comments
FCN-32s	59.0	Coarse segmentation; upsampling from last layer only
FCN-16s	62.5	Improved detail via skip connections
FCN-8s	64.0	Best performance; finest details captured

As supporting evidence to the quantitative results, the authors provide a figure showing the general architecture of the FCN model. The figure shows how an input image is subjected to several convolutional layers to obtain a feature map, and is subsequently pipelined through upsampling (or deconvolutional) layers coupled with skip connections to provide a pixel-wise segmentation map. Visualizing the setup is useful to communicate how the model recovers spatial information lost by downsampling and eventually produces more accurate semantic segmentation.

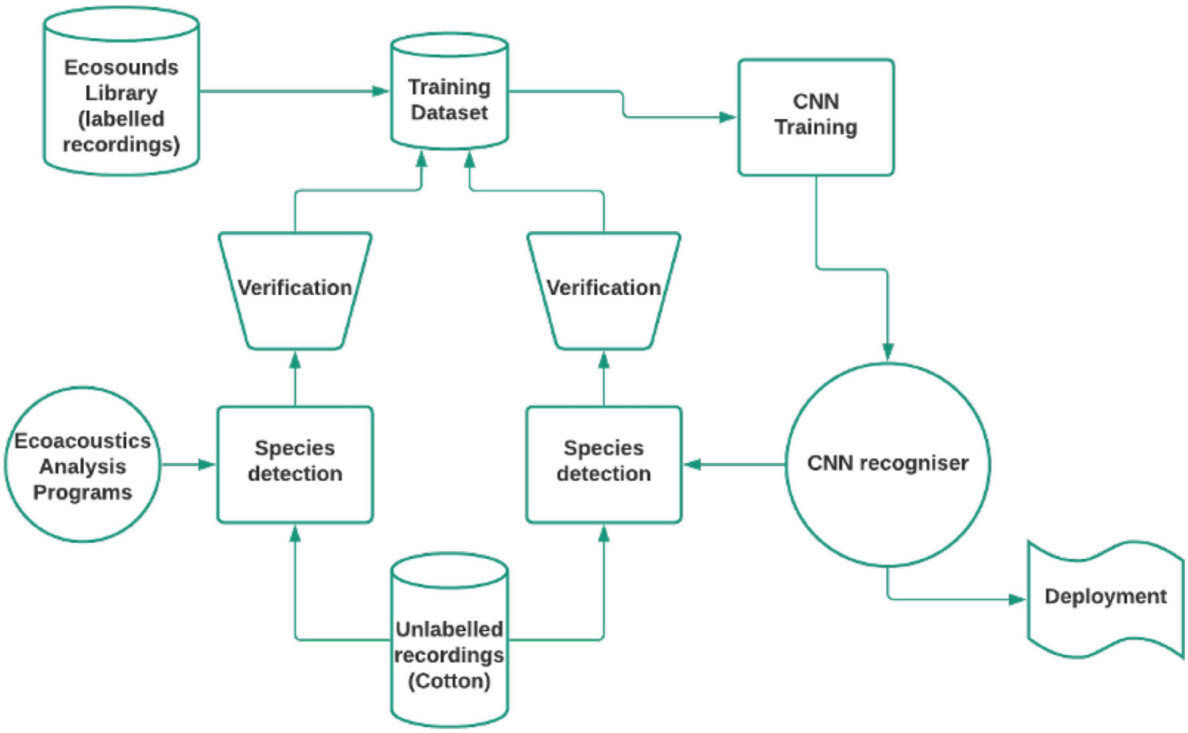

3. Contributions

The most important contribution of the present paper is to introduce a full convolutional architecture that converts traditional classification networks to segmentation networks. By

eliminating the fully connected layers and adding deconvolution layers, the approach enables end-to-end learning for dense prediction tasks. Another contribution is to use skip connections to combine multiple layers' features, enhancing segmentation performance by balancing coarse semantic and fine spatial context. In addition to simplifying the segmentation pipeline, it makes a dramatic gain in performance on hard benchmarks. All in all, the paper paved the way to all the subsequent advances in semantic segmentation and showed that deep, end-to-end trainable models could be used to achieve state-of-the-art performance without intrusive post-processing.

4. Criticism

I'm Despite all its novelty, FCN is not without constraints. One significant problem is that the operation of upsampling may result in blurred segmentation contours, as deconvolutional layers do not necessarily restore fine lost details in downsampling. Additionally, although skip connections make the model more accurate, they introduce difficulty in optimizing the network to perform optimally. The model will also not work with highly slender or tiny objects where fine contours are important. Finally, end-to-end training of these deep networks is computationally intensive and requires hyperparameters to be carefully optimized, something perhaps not so easy to accomplish for every user. These criticisms have led to subsequent studies working to find better architectures and post-processing techniques to overcome these drawbacks

5. Reference

Long, J., Shelhamer, E., & Darrell, T. (2015). Fully Convolutional Networks for Semantic Segmentation. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)